

Loan Data Exploration

July 29, 2020

0.1 Loan Data Analysis

```
[2]: #Import packages

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

# suppress warnings from final output
import warnings
warnings.simplefilter("ignore")
```

```
[3]: #load dataset into dataframe
df = pd.read_csv('prosperLoanData.csv')
```

```
[4]: print(df.shape)
print(df.dtypes)
```

```
(113937, 81)
ListingKey          object
ListingNumber       int64
ListingCreationDate object
CreditGrade        object
Term               int64
...
PercentFunded      float64
Recommendations    int64
InvestmentFromFriendsCount int64
InvestmentFromFriendsAmount float64
Investors          int64
Length: 81, dtype: object
```

```
[5]: df.head()
```

```
[5]:
```

	ListingKey	ListingNumber	ListingCreationDate	\
0	1021339766868145413AB3B	193129	2007-08-26 19:09:29.263000000	
1	10273602499503308B223C1	1209647	2014-02-27 08:28:07.900000000	
2	0EE9337825851032864889A	81716	2007-01-05 15:00:47.090000000	
3	0EF5356002482715299901A	658116	2012-10-22 11:02:35.010000000	
4	0F023589499656230C5E3E2	909464	2013-09-14 18:38:39.097000000	

	CreditGrade	Term	LoanStatus	ClosedDate	BorrowerAPR	\
0	C	36	Completed	2009-08-14 00:00:00	0.16516	
1	NaN	36	Current	NaN	0.12016	
2	HR	36	Completed	2009-12-17 00:00:00	0.28269	
3	NaN	36	Current	NaN	0.12528	
4	NaN	36	Current	NaN	0.24614	

	BorrowerRate	LenderYield	...	LP_ServiceFees	LP_CollectionFees	\
0	0.1580	0.1380	...	-133.18	0.0	
1	0.0920	0.0820	...	0.00	0.0	
2	0.2750	0.2400	...	-24.20	0.0	
3	0.0974	0.0874	...	-108.01	0.0	
4	0.2085	0.1985	...	-60.27	0.0	

	LP_GrossPrincipalLoss	LP_NetPrincipalLoss	LP_NonPrincipalRecoverypayments	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	PercentFunded	Recommendations	InvestmentFromFriendsCount	\
0	1.0	0	0	
1	1.0	0	0	
2	1.0	0	0	
3	1.0	0	0	
4	1.0	0	0	

	InvestmentFromFriendsAmount	Investors
0	0.0	258
1	0.0	1
2	0.0	41
3	0.0	158
4	0.0	20

[5 rows x 81 columns]

```
[6]: df.StatedMonthlyIncome.describe()
```

```
[6]: count    1.139370e+05
     mean     5.608026e+03
     std      7.478497e+03
     min      0.000000e+00
     25%      3.200333e+03
     50%      4.666667e+03
     75%      6.825000e+03
     max      1.750003e+06
     Name: StatedMonthlyIncome, dtype: float64
```

```
[7]: df.StatedMonthlyIncome.min()
```

```
[7]: 0.0
```

```
[8]: df.StatedMonthlyIncome.max()
```

```
[8]: 1750002.916667
```

```
[9]: df.MonthlyLoanPayment.describe()
```

```
[9]: count    113937.000000
     mean      272.475783
     std      192.697812
     min       0.000000
     25%      131.620000
     50%      217.740000
     75%      371.580000
     max      2251.510000
     Name: MonthlyLoanPayment, dtype: float64
```

```
[10]: df.LoanStatus.value_counts()
```

```
[10]: Current                56576
     Completed             38074
     Chargedoff            11992
     Defaulted              5018
     Past Due (1-15 days)    806
     Past Due (31-60 days)  363
     Past Due (61-90 days)  313
     Past Due (91-120 days) 304
     Past Due (16-30 days)  265
     FinalPaymentInProgress  205
     Past Due (>120 days)   16
     Cancelled               5
     Name: LoanStatus, dtype: int64
```

```
[11]: df.EmploymentStatus.value_counts()
```

```
[11]: Employed          67322
      Full-time       26355
      Self-employed    6134
      Not available    5347
      Other           3806
      Part-time       1088
      Not employed     835
      Retired          795
      Name: EmploymentStatus, dtype: int64
```

```
[12]: df[df.EmploymentStatus.isna()].count()
```

```
[12]: ListingKey          2255
      ListingNumber      2255
      ListingCreationDate 2255
      CreditGrade        2255
      Term               2255
      ...
      PercentFunded      2255
      Recommendations    2255
      InvestmentFromFriendsCount 2255
      InvestmentFromFriendsAmount 2255
      Investors          2255
      Length: 81, dtype: int64
```

```
[13]: df.EmploymentStatus = df.EmploymentStatus.fillna('Not available')
```

```
[14]: df.Occupation.value_counts()
```

```
[14]: Other              28617
      Professional      13628
      Computer Programmer  4478
      Executive         4311
      Teacher           3759
      ...
      Dentist           68
      Student - College Freshman  41
      Student - Community College  28
      Judge             22
      Student - Technical School  16
      Name: Occupation, Length: 67, dtype: int64
```

```
[15]: df.ListingCreationDate = pd.to_datetime(df.ListingCreationDate)
```

```
[16]: df.dtypes
```

```
[16]: ListingKey          object
      ListingNumber      int64
      ListingCreationDate datetime64[ns]
      CreditGrade         object
      Term               int64
      ...
      PercentFunded       float64
      Recommendations     int64
      InvestmentFromFriendsCount int64
      InvestmentFromFriendsAmount float64
      Investors           int64
      Length: 81, dtype: object
```

0.1.1 What is the structure of your dataset?

There are 113,937 different loans listed in this dataset with data on many factors regarding the loan (such as interest rate and length of the loan), as well as information on the borrower (such as occupation and employment status)

0.1.2 What is/are the main feature(s) of interest in your dataset?

It would be interesting to explore the loan status and see whether loan demographics or borrower demographics have a larger impact on whether a loan ends up defaulting. For this dataset, we will disregard late payments, but only look at loans that are either completed, defaulted or current on their payments.

0.1.3 What features in the dataset do you think will help support your investigation into your feature(s) of interest?


I suspect the borrower demographics (such as Prosper Score and the number of investors) will have a greater correlation with whether loans default than loan demographics (such as loan term length and interest rates for the borrower on the loan), so I will look at those features in this investigation.

0.2 Univariate Exploration

```
[17]: #All the people with no stated monthly income
      df[df['StatedMonthlyIncome'] == 0.0].StatedMonthlyIncome.count()
      df_noinc = df[df['StatedMonthlyIncome'] == 0.0]
```


```
[18]: plt.hist(data=df, x='StatedMonthlyIncome', bins=500);
      plt.xlabel('Stated Monthly Income')
      plt.ylabel('Count')
      plt.xlim(0,100000)
```

```
[18]: (0, 100000)
```



```
[19]: # there's a long tail in the distribution, so let's put it on a log scale
      ↪ instead
log_binsize = 0.05
bins = 10 ** np.arange(2.4, np.log10(df['StatedMonthlyIncome'].
      ↪ max())+log_binsize, log_binsize)


plt.figure(figsize=[8, 5])
plt.hist(data = df, x = 'StatedMonthlyIncome', bins = bins)
plt.xscale('log')
plt.xticks([500, 1e3, 2e3, 5e3, 1e4, 3e4], [500, '1k', '2k', '5k', '10k',
      ↪ '30k'])
plt.xlabel('Stated Monthly income ($) - log scale')
plt.ylabel('Count')
plt.title('Stated Monthly Income')
plt.show()
```



Stated Monthly Income for the borrower has a distribution with a long tail with a lot of stated monthly incomes on the lower end and few on the higher end. When plotted on a log scale, the distribution looks unimodal with a peak at around 4k on the log scale


```
[20]: col_order = ['Completed', 'Current', 'Defaulted']
      base_color = sns.color_palette()[0]
      plt.figure(figsize=[25, 10])
      sns.countplot(data = df, x = 'LoanStatus', color = base_color, order=col_order)
```

```
[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffa3fa15b50>
```



```
[21]: base_color = sns.color_palette()[0]
plt.figure(figsize=[12, 6])
sns.countplot(data = df, x = 'ProsperScore', color = base_color)
plt.title('Distribution of Prosper Score')
```

```
[21]: Text(0.5, 1.0, 'Distribution of Prosper Score')
```




```
[22]: # let's plot all three together to get an idea of each ordinal variable's
      ↳ distribution.
col_order = ['Completed', 'Current', 'Defaulted']

fig, ax = plt.subplots(nrows=3, figsize = [25,12])

default_color = sns.color_palette()[0]
sns.countplot(data = df, x = 'ProsperScore', color = default_color, ax = ax[0])
sns.countplot(data = df, x = 'Term', color = default_color, ax = ax[1])
sns.countplot(data = df, x = 'LoanStatus', color = default_color,
      ↳ order=col_order, ax = ax[2])

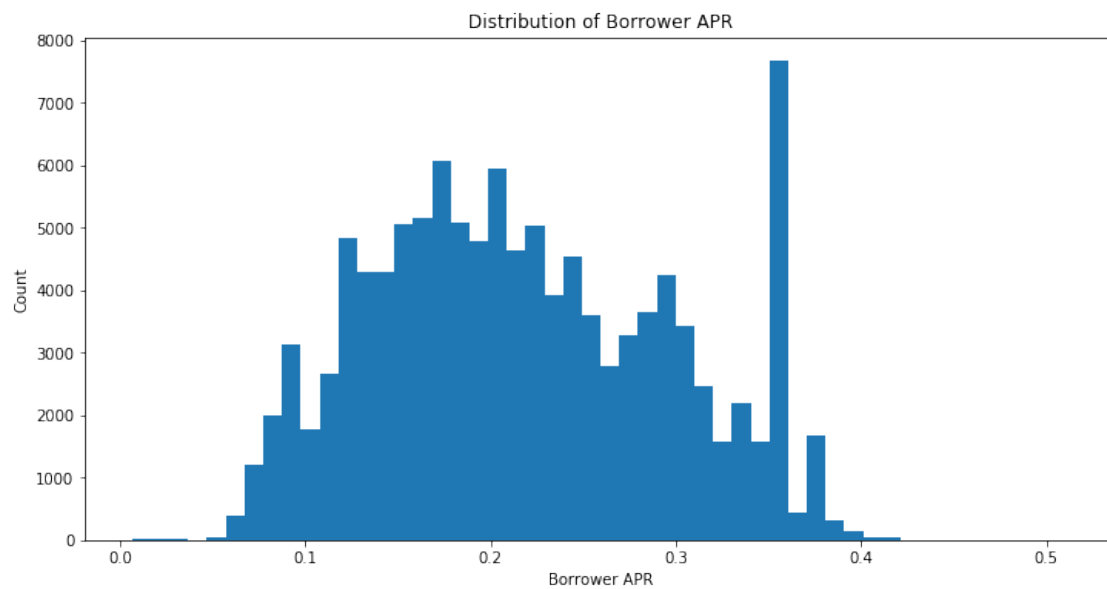
for ax in fig.axes:
    ax.tick_params(labelrotation=15)

plt.show()
```

It looks like a majority of loans have a term of 36 months and borrowers tend to have a salary in the \$25-75k range

```
[48]: #plot of Borrower APR distribution
plt.figure(figsize=[12, 6])
plt.hist(data=df, x='BorrowerAPR', bins=50);
plt.xlabel('Borrower APR')
plt.ylabel('Count')
plt.title('Distribution of Borrower APR');
```



```
[26]: # convert IncomeRange into ordered categorical types
ordinal_var_dict = {'IncomeRange': ['Not displayed', 'Not_
    ↳employed', '$0', '$1-24,999', '$25,000-49,999',
    ↳
    ↳'$50,000-74,999', '$75,000-99,999', '$100,000+'],
    ↳
    ↳'LoanStatus' :_
    ↳['Current', 'Completed', 'Chargedoff', 'Cancelled', 'Defaulted', 'FinalPaymentInProgress',
    ↳
    ↳
    ↳'Past Due (1-15 days)', 'Past Due (16-30_
    ↳days)', 'Past Due (31-60 days)',
    ↳
    ↳
    ↳'Past Due (61-90 days)', 'Past Due (91-120_
    ↳days)', 'Past Due (>120 days)']
    }

for var in ordinal_var_dict:
    ordered_var = pd.api.types.CategoricalDtype(ordered = True,
    ↳
    ↳categories =_
    ↳ordinal_var_dict[var])
    df[var] = df[var].astype(ordered_var)
```

```
[27]: df['IncomeRange'] = df.IncomeRange.astype('category')
df['LoanStatus'] = df.LoanStatus.astype('category')
```

```
[28]: df.IncomeRange.dtypes
```

```
[28]: CategoricalDtype(categories=['Not displayed', 'Not employed', '$0', '$1-24,999',
    '$25,000-49,999', '$50,000-74,999', '$75,000-99,999',
    '$100,000+'],
    ordered=True)
```


```
[29]: df.IncomeRange.value_counts()
```

```
[29]: $25,000-49,999    32192
$50,000-74,999    31050
$100,000+        17337
$75,000-99,999    16916
Not displayed      7741
$1-24,999         7274
Not employed       806
$0                621
Name: IncomeRange, dtype: int64
```

```
[30]: inc_order =_
    ↳['$0', '$1-24,999', '$25,000-49,999', '$50,000-74,999', '$75,000-99,999', '$100,000+']
base_color = sns.color_palette()[0]
plt.figure(figsize=[12, 6])
sns.countplot(data = df, x = 'IncomeRange', color = base_color, order=inc_order)
```

```
plt.title('Income Brackets for Loan Applicants')
```

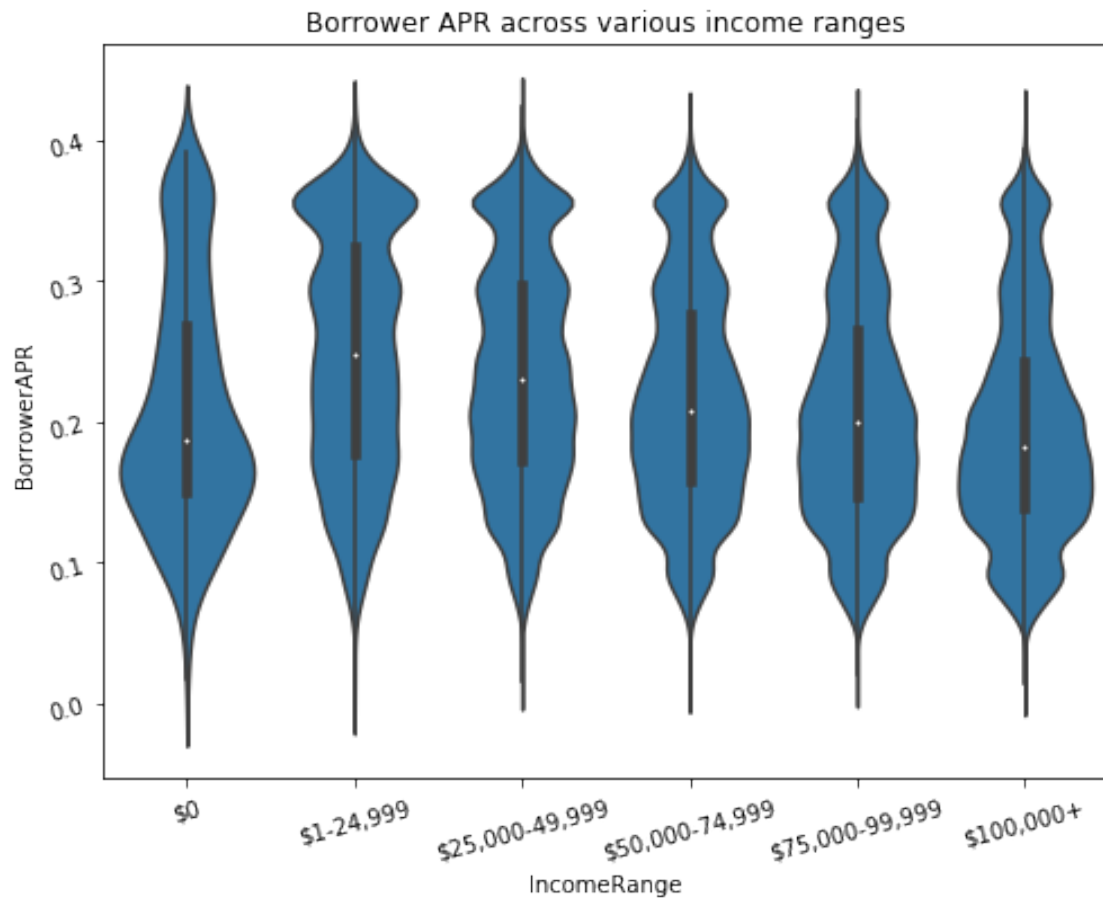
```
[30]: Text(0.5, 1.0, 'Income Brackets for Loan Applicants')
```



```
[31]: df.IncomeRange.value_counts()
```


```
[31]: $25,000-49,999    32192
      $50,000-74,999    31050
      $100,000+         17337
      $75,000-99,999    16916
      Not displayed      7741
      $1-24,999          7274
      Not employed        806
      $0                  621
      Name: IncomeRange, dtype: int64
```

```
[32]: #plot income of the borrower vs. the APR that the borrower receives
inc_order = [
    '$0', '$1-24,999', '$25,000-49,999', '$50,000-74,999', '$75,000-99,999', '$100,000+'
]
base_color = sns.color_palette()[0]
plt.figure(figsize = [8, 6])
sns.violinplot(data=df, x='IncomeRange', y='BorrowerAPR', order=inc_order,
    color=base_color);
plt.tick_params(labelrotation=15)
plt.title('Borrower APR across various income ranges');
```



```
[49]: #plot income of the borrower vs. the APR that the borrower receives
col_order = ['Completed', 'Current', 'Defaulted']

base_color = sns.color_palette()[0]
plt.figure(figsize = [8, 6])
sns.violinplot(data=df, x='LoanStatus', y='BorrowerAPR', order=col_order,
               color=base_color);
plt.tick_params(labelrotation=15)
plt.title('Borrower APR across loan status');
```



0.2.1 Discuss the distribution(s) of your variable(s) of interest. Were there any unusual points? Did you need to perform any transformations?

It is interesting to note that the shape of the distribution of ProsperScores is seemingly symmetric, and that there is less ratings in the 5 or 7 category.

Also, we see that the distribution of Stated Monthly Income of the borrower is extremely right skewed due to the fact that there are a few borrowers that have an extremely large income in comparison to most borrowers. So, in order to clearly visualize this data, we need to apply a log transform on the distribution. Once the transform is applied, we see that the data is unimodal.

0.2.2 Of the features you investigated, were there any unusual distributions? Did you perform any operations on the data to tidy, adjust, or change the form of the data? If so, why did you do this?

Some of the categorical variables were not ordered in a very clear way, including the income range for the borrower as well as the status of the loan, so I organized them a clearer way based on ordinality.

I took a log transform of the Stated Monthly Income to visualize it better at it had a long tail with a few outliers (people who had a very large income in comparison to the rest that skewed the data


and made it difficult to see trends)

0.3 Bivariate Exploration

Start off with looking at Pairwise correlations of the data

```
[33]: numeric_vars = ['MonthlyLoanPayment', 'Recommendations', 'StatedMonthlyIncome', 'InvestmentFromFriendsAmount', 'Investors', 'BorrowerAPR']
categoric_vars = ['BorrowerState', 'LoanStatus', 'IncomeRange', 'Term']
```

```
[34]: # correlation plot
plt.figure(figsize = [8, 5])
sns.heatmap(df[numeric_vars].corr(), annot = True, fmt = '.3f',
            cmap = 'vlag_r', center = 0)
plt.show()
```




Something interesting to note from this heatmap is that the strongest correlation is the relationship between investors and monthly loan payment, and that it is positive, meaning that these may be related with one another.

Also, we see that there is very little correlation between the stated monthly income of the borrower and the investment that they receive from friends.

```
[39]: col_order = ['Completed', 'Current', 'Defaulted']
base_color = sns.color_palette()[0]
plt.figure(figsize=[25, 5])
sns.stripplot(x="LoanStatus", y="ProsperScore", data=df, color=base_color,
→order=col_order, jitter=True, dodge=True)
```

```
[39]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffa402e7f90>
```



```
[40]: df[df['LoanStatus'] == 'Defaulted'].ProsperScore.mean()
```


```
[40]: 5.619900497512438
```

```
[41]: df[df['LoanStatus'] == 'Completed'].ProsperScore.mean()
```

```
[41]: 6.536513425549227
```


0.4 Multivariate Exploration

```
[42]: col_order = ['Completed', 'Current', 'Defaulted']
g = sns.FacetGrid(data = df, col = 'LoanStatus', col_order = col_order)
g.map(sns.pointplot, 'ProsperScore', 'BorrowerAPR', linestyle = '', order = [1.
→0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0])
plt.show()
```



We see here that when we compare loans that defaulted with current loans, we see that Borrower APR tended to be higher regardless of the Prosper Score of the Borrower. This is interesting to note as we try to determine what factors may be contributing to whether or not a loan defaults.


```
[43]: col_order = ['Completed', 'Current', 'Defaulted']
fig_dims = (25, 5)
fig, ax = plt.subplots(figsize=fig_dims)
sns.barplot(x="LoanStatus", y="BorrowerAPR", hue="Term", data=df,
            order=col_order, ax=ax);
```



Something interesting to note from this plot is that regardless of the term of the loan, Borrower APR tended to be higher for loans that defaulted vs. those that were current or completed, though it is seen most significantly among the short term (12 month long) loans, where the APR for defaulted loans is significantly higher.

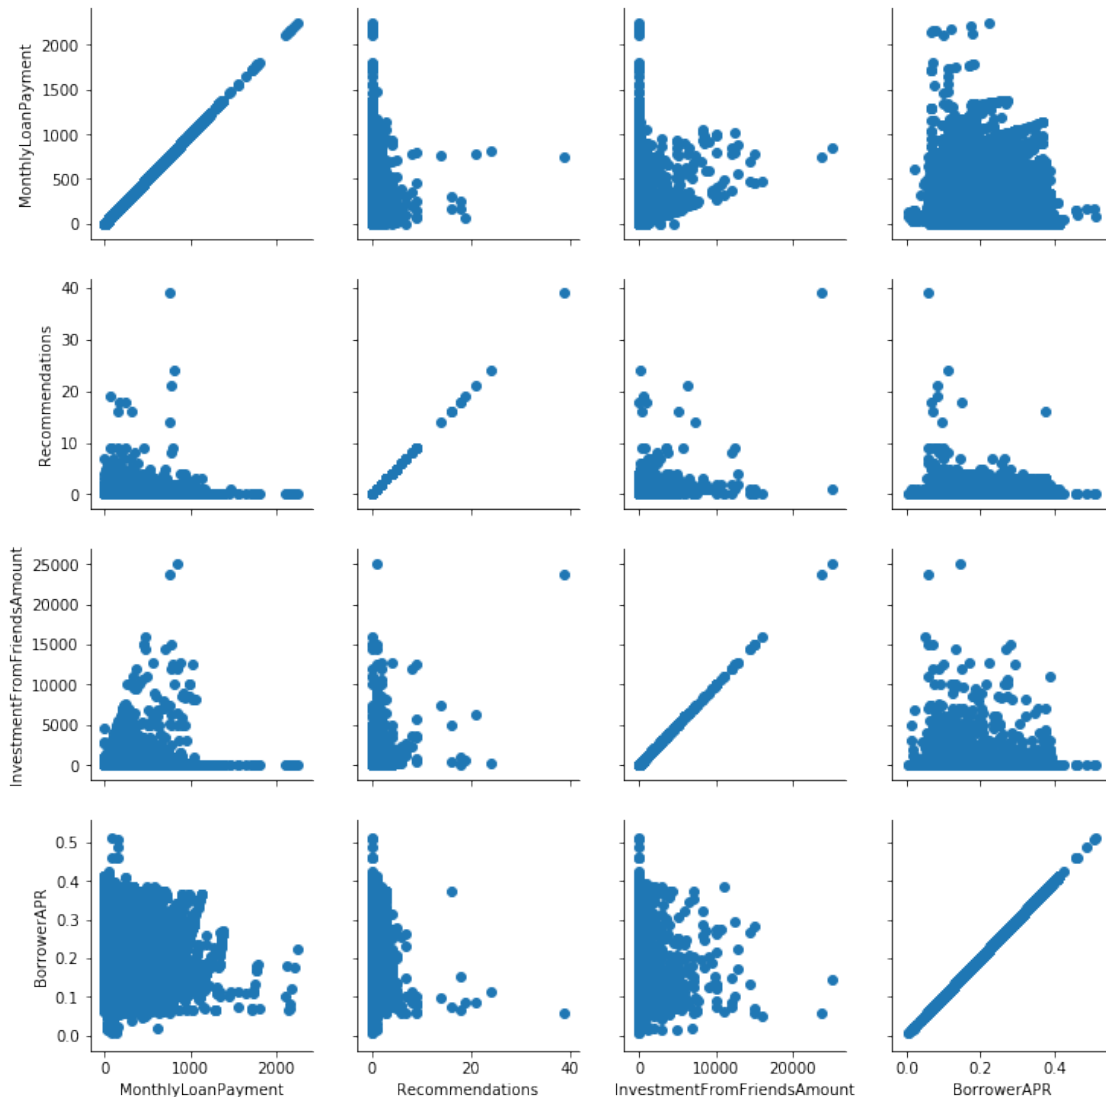
This is different from the results obtained when we looked at the relationship between loan status and Borrower APR only.

```
[44]: col_order = ['Completed', 'Current', 'Defaulted']
fig_dims = (25, 5)
fig, ax = plt.subplots(figsize=fig_dims)
sns.barplot(x="LoanStatus", y="ProsperScore", hue="Term", data=df,
            order=col_order, ax=ax);
```



When looking at ProsperScore, however, we see a very interesting phenomenon. For 60 month loans, the Prosper Score tends to be pretty similar regardless of whether a loan defaults or not. So it doesn't seem like it would be a good indicator to determine if a loan will default in the long term, though for short term loans, will still see the Prosper Score tends to be lower for loans that default.


```
[45]: numeric_vars = ['MonthlyLoanPayment', 'Recommendations',  
    ↪ 'InvestmentFromFriendsAmount', 'BorrowerAPR']  
categoric_vars = ['BorrowerState', 'LoanStatus', 'IncomeRange', 'Term']  
  
g= sns.PairGrid(data=df, vars=numeric_vars)  
g.map(plt.scatter);
```



Something interesting to see in this plot is that it seems that there is a slightly negative correlation between the number of recommendations that a borrower receives and the amount of investment that they receive from friends. This seems counterintuitive as we would expect that people would want to invest more in loans that have more recommendations, but this is not the case.

Factors by State Something else of interest is to see if the location of the borrower has any affect on if a loan defaults. Since there are several states to look at and it would difficult to visualize them all, we will look only at the top 4 states with the most amount of loans.

```
[46]: df.BorrowerState.value_counts().head()
```


```
[46]: CA      14717
      TX      6842
      NY      6729
      FL      6720
      IL      5921
      Name: BorrowerState, dtype: int64
```

According to the above, the four states with the most loans are California, Texas, New York, and Florida. So we will look at loans in those states.

```
[47]: states = ['CA', 'TX', 'NY', 'FL']
      col_order = ['Completed', 'Current', 'Defaulted']
      fig_dims = (25, 5)
      fig, ax = plt.subplots(figsize=fig_dims)

      #default_color = sns.color_palette()[0]
      sns.barplot(data = df, x = 'BorrowerState', y = 'BorrowerAPR',
                  hue='LoanStatus', order=states, hue_order=col_order, ax=ax)
```

```
[47]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffa3f8ad050>
```



0.4.1 Talk about some of the relationships you observed in this part of the investigation. Were there features that strengthened each other in terms of looking at your feature(s) of interest?

Of all the factors observed throughout this investigation, the one that seemed to have the most effect on whether a loan defaulted or not was the term of the loan, though these differences may be more pronounced because of the fact that there are not nearly as many 12 month loans as compared to the longer ones

0.4.2 Were there any interesting or surprising interactions between features?

There was an interesting anomaly that occurred when loans were broken up by term within different loan statuses. The 12 month loans for defaulted loans had an unusually high Borrower APR (on average) in comparison to Completed and Current loans of the same term. This is not very clearly seen when just looking at loan status and APR by themselves.

[]: